

**Course:** ENSF 611 – Fall 2022

**Lab #:** 0

**Instructor:** Kumar

**Student Name:** Ardit Baboci

**Submission Date:** September 26, 2022

# lab0-pandas-auto\_mpg

September 25, 2022

## 1 Pandas

As described at <https://pandas.pydata.org> > pandas is an open source, BSD-licensed library providing high-performance, easy-to-use data structures and data analysis tools for the Python programming language.

### 1.1 Resources

1. Ch 5-6 in Python for Data Analysis, 2nd Ed, Wes McKinney (Ucalgary library and <https://github.com/wesm/pydata-book>)
2. Ch 3 in Python Data Science Handbook, Jake VanderPlas (Ucalgary library and <https://github.com/jakevdp/PythonDataScienceHandbook>)

Let's explore some of the features.

First, import Pandas, and Numpy as a good companion.

```
[1]: import numpy as np
import pandas as pd
```

### 1.2 Create pandas DataFrames

There are several ways to create Pandas DataFrames, most notably from reading a csv (comma separated values file). DataFrames are 'spreadsheets' in Python. We will often use `df` as a variable name for a DataFrame.

If data is not stored in a file, a DataFrame can be created from a dictionary of lists

```
data = {'state': ['Ohio', 'Ohio', 'Ohio', 'Nevada', 'Nevada', 'Nevada'],
        'year': [2000, 2001, 2002, 2001, 2002, 2003],
        'pop': [1.5, 1.7, 3.6, 2.4, 2.9, 3.2]}
frame = pd.DataFrame(data)
```

where dictionary keys become column headers.

An alternative is to create from a numpy array and set column headers separately:

```
[2]: # From a numpy array
df = pd.DataFrame(np.arange(20).reshape(5,4), columns=['alpha', 'beta', 'gamma', 'delta'])
df
```

```
[2]:
```

	alpha	beta	gamma	delta
0	0	1	2	3
1	4	5	6	7
2	8	9	10	11
3	12	13	14	15
4	16	17	18	19

```
[3]: # checking its type
type(df)
```

```
[3]: pandas.core.frame.DataFrame
```

### 1.3 Indexing

Accessing data in Dataframes is done by rows and columns, either index or label based.

```
[4]: # select a column
df['alpha']
```

```
[4]:
```

0	0
1	4
2	8
3	12
4	16

Name: alpha, dtype: int32

```
[5]: # select two columns
df[['alpha', 'gamma']]
```

```
[5]:
```

	alpha	gamma
0	0	2
1	4	6
2	8	10
3	12	14
4	16	18

```
[6]: # select rows
df.iloc[:2]
```

```
[6]:
```

	alpha	beta	gamma	delta
0	0	1	2	3
1	4	5	6	7

```
[7]: # select rows and columns
df.iloc[:2, :2]
```

```
[7]:
```

	alpha	beta
0	0	1

1        4        5

```
[8]: # select rows and columns, mixed
df.loc[:2, ['alpha', 'beta']]
```

```
[8]:     alpha  beta
0        0     1
1        4     5
2        8     9
```

## 1.4 DataFrame math

Similar to Numpy, DataFrames support direct math

```
[9]: # direct math
df2 = (9/5) * df + 32
df2
```

```
[9]:     alpha  beta  gamma  delta
0     32.0  33.8   35.6   37.4
1     39.2  41.0   42.8   44.6
2     46.4  48.2   50.0   51.8
3     53.6  55.4   57.2   59.0
4     60.8  62.6   64.4   66.2
```

```
[10]: # add two dataframes of same shape
df + df2
```

```
[10]:     alpha  beta  gamma  delta
0     32.0  34.8   37.6   40.4
1     43.2  46.0   48.8   51.6
2     54.4  57.2   60.0   62.8
3     65.6  68.4   71.2   74.0
4     76.8  79.6   82.4   85.2
```

```
[11]: # map a function to each column
f = lambda x: x.max() - x.min()

df.apply(f)
```

```
[11]: alpha     16
beta     16
gamma     16
delta     16
dtype: int64
```

## 1.5 DataFrame manipulation

Adding and deleting columns, as well as changing entries is similar to Python dictionaries.

Note that most DataFrame methods do not change the DataFrame directly, but return a new DataFrame. It is always good to check how the method you are invoking behaves.

```
[12]: # add a column
df['epsilon'] = ['low', 'medium', 'low', 'high', 'high']
df
```

```
[12]:   alpha  beta  gamma  delta  epsilon
0      0     1      2      3      low
1      4     5      6      7  medium
2      8     9     10     11      low
3     12    13     14     15     high
4     16    17     18     19     high
```

```
[13]: # What is the size?
df.shape
```

```
[13]: (5, 5)
```

```
[15]: # delete column
df_dropped = df.drop(columns=['gamma'])
df_dropped
```

```
[15]:   alpha  beta  delta  epsilon
0      0     1      3      low
1      4     5      7  medium
2      8     9     11      low
3     12    13     15     high
4     16    17     19     high
```

```
[16]: # the original dataframe is unaffected
df
```

```
[16]:   alpha  beta  gamma  delta  epsilon
0      0     1      2      3      low
1      4     5      6      7  medium
2      8     9     10     11      low
3     12    13     14     15     high
4     16    17     18     19     high
```

Let's create a copy and assign new values to the first column:

```
[17]: df_copy = df.copy()
df_copy['alpha'] = 20
print(df)
print(df_copy)
```

```
   alpha  beta  gamma  delta  epsilon
0      0     1      2      3      low
```

1	4	5	6	7	medium
2	8	9	10	11	low
3	12	13	14	15	high
4	16	17	18	19	high
	alpha	beta	gamma	delta	epsilon
0	20	1	2	3	low
1	20	5	6	7	medium
2	20	9	10	11	low
3	20	13	14	15	high
4	20	17	18	19	high

DataFrames can be sorted by column:

```
[18]: # sorting values
df.sort_values(by='epsilon')
```

```
[18]:   alpha  beta  gamma  delta  epsilon
3     12    13     14     15     high
4     16    17     18     19     high
0      0     1      2      3      low
2      8     9     10     11     low
1      4     5      6      7    medium
```

## 1.6 Load data from file

Most often data will come from somewhere, often csv files, and using `pd.read_csv()` will allow smooth creation of DataFrames.

Let's load that same heart-attack.csv that we used in Numpy before:

```
[2]: data = pd.read_csv('auto-mpg.data.csv')
```

After loading data, it is good practice to check what we have. Usually, the sequences is: 1. Check dimension 2. Peek at the first rows 3. Get info on data types and missing values 4. Summarize columns

```
[3]: # Check dimension (rows, columns)
data.shape
```

```
[3]: (398, 9)
```

```
[4]: # Peek at the first rows
data.head()
```

```
[4]:   mpg  cylinders  displacement  horsepower  weight  acceleration  model year  \
0  18.0          8          307.0          130   3504           12.0         70
1  15.0          8          350.0          165   3693           11.5         70
2  18.0          8          318.0          150   3436           11.0         70
3  16.0          8          304.0          150   3433           12.0         70
4  17.0          8          302.0          140   3449           10.5         70
```

```

      origin      car name
0         1  chevrolet chevelle malibu
1         1      buick skylark 320
2         1    plymouth satellite
3         1      amc rebel sst
4         1      ford torino

```

```
[5]: # Column names are
      data.columns
```

```
[5]: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
           'acceleration', 'model year', 'origin', 'car name'],
          dtype='object')
```

```
[6]: # Get info on data types and missing values
      data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg              398 non-null   float64
1   cylinders        398 non-null   int64
2   displacement     398 non-null   float64
3   horsepower       398 non-null   object
4   weight           398 non-null   int64
5   acceleration     398 non-null   float64
6   model year      398 non-null   int64
7   origin           398 non-null   int64
8   car name        398 non-null   object
dtypes: float64(3), int64(4), object(2)
memory usage: 28.1+ KB

```

## 1.7 Summarize values

What is the mean, std, min, max in each column?

```
[7]: data.mean()
```

```

C:\Users\Ardit\AppData\Local\Temp\ipykernel_11480\531903386.py:1: FutureWarning:
Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None')
is deprecated; in a future version this will raise TypeError.  Select only valid
columns before calling the reduction.
      data.mean()

```

```
[7]: mpg          23.514573
      cylinders    5.454774
```

```

displacement    193.425879
weight          2970.424623
acceleration     15.568090
model year      76.010050
origin          1.572864
dtype: float64

```

```

[8]: # where are the other columns? Check data types
data.dtypes

```

```

[8]: mpg          float64
cylinders        int64
displacement     float64
horsepower       object
weight          int64
acceleration     float64
model year       int64
origin           int64
car name         object
dtype: object

```

Notice that many columns are of type object, which is not a number. Maybe this has to do with missing values? We know from peeking at the first rows that there are '?' values in there. Let's replace these with the string NaN for not-a-number.

```

[9]: # replace '?' with 'NaN'
data = data.replace({'?': 'NaN'})
data.head()

```

```

[9]:      mpg  cylinders  displacement  horsepower  weight  acceleration  model year  \
0  18.0         8         307.0         130    3504         12.0         70
1  15.0         8         350.0         165    3693         11.5         70
2  18.0         8         318.0         150    3436         11.0         70
3  16.0         8         304.0         150    3433         12.0         70
4  17.0         8         302.0         140    3449         10.5         70

      origin      car name
0         1  chevrolet chevelle malibu
1         1      buick skylark 320
2         1    plymouth satellite
3         1      amc rebel sst
4         1      ford torino

```

Pandas knows that 'NaN' probably means that numbers are missing. Now we can convert the data type from object to float

```

[10]: # convert dtypes
data.iloc[:, :8] = data.iloc[:, :8].astype('float')

```



```
data.dtypes
```

```
[10]: mpg          float64
      cylinders    float64
      displacement float64
      horsepower   float64
      weight       float64
      acceleration float64
      model year   float64
      origin       float64
      car name     object
      dtype: object
```

We could have loaded the data with the `na_values` argument to indicate that ‘?’ means missing number:

```
[11]: data = pd.read_csv('auto-mpg.data.csv', na_values='?')
      data.dtypes
```

```
[11]: mpg          float64
      cylinders     int64
      displacement float64
      horsepower    float64
      weight        int64
      acceleration  float64
      model year    int64
      origin        int64
      car name      object
      dtype: object
```

This worked nicely. Now we can describe all columns, meaning printing basic statistics. Note that by default Pandas ignores NaN, whereas Numpy does not.

```
[12]: data.describe() # ignores NaN
```

```
[12]:
```

	mpg	cylinders	displacement	horsepower	weight	\
count	398.000000	398.000000	398.000000	392.000000	398.000000	
mean	23.514573	5.454774	193.425879	104.469388	2970.424623	
std	7.815984	1.701004	104.269838	38.491160	846.841774	
min	9.000000	3.000000	68.000000	46.000000	1613.000000	
25%	17.500000	4.000000	104.250000	75.000000	2223.750000	
50%	23.000000	4.000000	148.500000	93.500000	2803.500000	
75%	29.000000	8.000000	262.000000	126.000000	3608.000000	
max	46.600000	8.000000	455.000000	230.000000	5140.000000	

	acceleration	model year	origin
count	398.000000	398.000000	398.000000
mean	15.568090	76.010050	1.572864

std	2.757689	3.697627	0.802055
min	8.000000	70.000000	1.000000
25%	13.825000	73.000000	1.000000
50%	15.500000	76.000000	1.000000
75%	17.175000	79.000000	2.000000
max	24.800000	82.000000	3.000000

We could be interested by these statistics in each of the origins. To get these, we first group values by origin, then ask for the description. We will only look at mpg for clarity

```
[13]: data.groupby(by='origin').describe().mpg
```

```
[13]:
```

	count	mean	std	min	25%	50%	75%	max
origin								
1	249.0	20.083534	6.402892	9.0	15.0	18.5	24.00	39.0
2	70.0	27.891429	6.723930	16.2	24.0	26.5	30.65	44.3
3	79.0	30.450633	6.090048	18.0	25.7	31.6	34.05	46.6

## 1.8 Find NaNs

How many NaNs in each column?

We can ask which entries are null, which produces a boolean array

```
[14]: data.isnull()
```

```
[14]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	\
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	
..	...	...	...	...	...	...	
393	False	False	False	False	False	False	
394	False	False	False	False	False	False	
395	False	False	False	False	False	False	
396	False	False	False	False	False	False	
397	False	False	False	False	False	False	

	model	year	origin	car name
0	False	False	False	
1	False	False	False	
2	False	False	False	
3	False	False	False	
4	False	False	False	
..	...	...	...	
393	False	False	False	
394	False	False	False	
395	False	False	False	

```

396      False  False  False
397      False  False  False

```

```
[398 rows x 9 columns]
```

Applying `sum()` to this boolean array will count the number of `True` values in each column

```
[15]: data.isnull().sum()
```

```

[15]: mpg          0
      cylinders    0
      displacement  0
      horsepower    6
      weight        0
      acceleration  0
      model year    0
      origin        0
      car name      0
      dtype: int64

```

We get complementary information from `info()`

```
[16]: data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   mpg             398 non-null   float64
 1   cylinders        398 non-null   int64
 2   displacement     398 non-null   float64
 3   horsepower       392 non-null   float64
 4   weight           398 non-null   int64
 5   acceleration     398 non-null   float64
 6   model year       398 non-null   int64
 7   origin           398 non-null   int64
 8   car name         398 non-null   object
dtypes: float64(4), int64(4), object(1)
memory usage: 28.1+ KB

```

We can fill (replace) these missing values, for example with the minimum value in each column

```
[17]: data.fillna(data.min()).describe()
```

```

[17]:          mpg  cylinders  displacement  horsepower  weight  \
count  398.000000  398.000000    398.000000    398.000000  398.000000
mean    23.514573    5.454774    193.425879    103.587940  2970.424623
std      7.815984    1.701004    104.269838     38.859575   846.841774

```

min	9.000000	3.000000	68.000000	46.000000	1613.000000
25%	17.500000	4.000000	104.250000	75.000000	2223.750000
50%	23.000000	4.000000	148.500000	92.000000	2803.500000
75%	29.000000	8.000000	262.000000	125.000000	3608.000000
max	46.600000	8.000000	455.000000	230.000000	5140.000000

	acceleration	model year	origin
count	398.000000	398.000000	398.000000
mean	15.568090	76.010050	1.572864
std	2.757689	3.697627	0.802055
min	8.000000	70.000000	1.000000
25%	13.825000	73.000000	1.000000
50%	15.500000	76.000000	1.000000
75%	17.175000	79.000000	2.000000
max	24.800000	82.000000	3.000000

## 1.9 Count unique values (a histogram)

We finish off, with our good friend the histogram

```
[18]: data['mpg'].value_counts()
```

```
[18]: 13.0    20
      14.0    19
      18.0    17
      15.0    16
      26.0    14
      ..
      31.9     1
      16.9     1
      18.2     1
      22.3     1
      44.0     1
      Name: mpg, Length: 129, dtype: int64
```

```
[ ]:
```

# lab0-visualization-auto\_mpg

September 25, 2022

## 1 Visualization

### 1.1 Topics

1. Matplotlib core framework
2. Pandas plot()
3. Seaborn statistical visualization
4. (not covered) Grammar of graphics (ggplot2 see plotnine )
5. (not covered) Interactive plotting

### 1.2 Resources

1. Ch 9 in Python for Data Analysis, 2nd Ed, Wes McKinney (Ucalgary library and <https://github.com/wesm/pydata-book>)
2. Ch 4 in Python Data Science Handbook, Jake VanderPlas (Ucalgary library and <https://github.com/jakevdp/PythonDataScienceHandbook>)
3. Fundamentals of Data Visualization, Claus O. Wilke (Ucalgary library and <https://serialmentor.com/dataviz/index.html>)
4. Overview by Jake VanderPlas <https://www.youtube.com/watch?v=FytuB8nFHPQ>

### 1.3 Matplotlib

Matplotlib is a Python 2D plotting library which produces publication quality figures in a variety of hardcopy formats and interactive environments across platforms.

Matplotlib tries to make easy things easy and hard things possible.

For simple plotting the pyplot module provides a MATLAB-like interface

<https://matplotlib.org>

Importing matplotlib looks like this

```
[1]: %matplotlib inline

import numpy as np
import pandas as pd
import matplotlib as mpl
import matplotlib.pyplot as plt
```

### 1.3.1 Two interfaces

There are two ways to interact with Matplotlib: a Matlab style and an object oriented style interface.


See Ch 4 in Python Data Science Handbook, Jake VanderPlas

- Two Interfaces for the Price of One, pp. 222
- Matplotlib Gotchas, pp. 232

### 1.3.2 Matlab style interface

```
[2]: x = np.linspace(0, 10, 100)

plt.plot(x, np.sin(x), '-')
plt.plot(x, np.cos(x), '--');
```




Adding decorations to the plot is done by repeatedly calling functions on the imported `plt` module. All calls within the cell will be applied to the current figure and axes.

```
[3]: plt.plot(x, np.sin(x), '-', label='sine')
plt.plot(x, np.cos(x), '--', label='cosine')

plt.xlim([0, 10])
plt.ylim([-1, 1])
plt.xlabel('x')
plt.ylabel('amplitude')
plt.title('My sines')
```

```
plt.grid()  
plt.legend();
```




### 1.3.3 Object oriented interface

With this interface, you first create a figure and an axes object, then call their methods to change the plot.

```
[4]: fig = plt.figure()  
     ax = plt.axes()  
     ax.plot(x, np.sin(x), '-')
```


```
     ax.plot(x, np.cos(x), '--');
```



```
[5]: fig = plt.figure()
ax = plt.axes()
ax.plot(x, np.sin(x), '-', label='sine')
ax.plot(x, np.cos(x), '--', label='cosine')

ax.set(xlim=[0, 10], ylim=[-1, 1],
       xlabel='x', ylabel='amplitude',
       title='My sines');
ax.grid()
ax.legend();
```





### 1.3.4 Save to file

With the figure object at hand, we can save to file

```
[6]: fig.savefig('sines.pdf')
```

```
[7]: ls *.pdf
```

Volume in drive C has no label.

Volume Serial Number is 94FF-36B3

Directory of C:\Users\Ardit\OneDrive - University of Calgary\ENSF  
611\ensf611-labs\lab0

2022-09-25	07:29 PM	(73,893)	Lab0_ababoci.pdf
2022-09-25	07:26 PM	66,030	lab0-pandas-auto_mpg.pdf
2022-09-25	07:32 PM	13,528	sines.pdf
	3 File(s)	153,451 bytes	
	0 Dir(s)	745,490,845,696 bytes free	

## 1.4 Plotting with pandas

We use the standard convention for referencing the matplotlib API ... We provide the basics in pandas to easily create decent looking plots.

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/visualization.html](https://pandas.pydata.org/pandas-docs/stable/user_guide/visualization.html)


Let's load the auto-mpg dataset

```
[6]: data = pd.read_csv('auto-mpg.data.csv', na_values='?')
```

Plotting all columns, works, but does not provide a lot of insight.


```
[7]: data.plot()
```

```
[7]: <matplotlib.axes._subplots.AxesSubplot at 0x1631df5d288>
```



Let's look at the mpg distribution (a histogram)

```
[8]: data['mpg'].plot.hist();
```



How many samples do we have in each origin?

```
[9]: data.origin.value_counts()
```

```
[9]: 1    249
      3     79
      2     70
      Name: origin, dtype: int64
```


Notice that we accessed the origin column with dot notation. This can be done whenever the column name is 'nice' enough to be a python variable name.

Do we have similar mpgs in each origin?

Plotting three histograms for each origin side beside directly form the dataframe:


```
[10]: axs = data.hist(column='mpg', by='origin')
      print(axs)
```

```
[[<matplotlib.axes._subplots.AxesSubplot object at 0x000001631F0EAE88>
  <matplotlib.axes._subplots.AxesSubplot object at 0x000001631F13B4C8>]
 [<matplotlib.axes._subplots.AxesSubplot object at 0x000001631F16D708>
  <matplotlib.axes._subplots.AxesSubplot object at 0x000001631F1A2AC8>]]
```




To format this plot, we can work on the axes (array) that is returned by the plot call. We use Matplotlib object oriented interface methods to do this

```
[11]: axs = data.hist(column='mpg', by='origin')
axs[0][0].set(title='Origin 1', ylim=[0, 45])
axs[0][1].set(title='Origin 2', ylim=[0, 45])
axs[1][0].set(title='Origin 3', ylim=[0, 45]);
```



Is mpg and model year correlated? Maybe it is different for the different origins?  
Let's have a look with a scatter plot.

```
[12]: data.plot.scatter('mpg', 'model year', c='origin', colormap='winter');
```




According to:

<https://stackoverflow.com/questions/43578976/pandas-missing-x-tick-labels>

the missing x-labels are a pandas bug.

Workaround is to create axes prior to calling plot

```
[13]: fig, ax = plt.subplots()
      data.plot.scatter('mpg', 'model year', c='origin', colormap='viridis', ax=ax);
```




It is a bit annoying that there is a colorbar, we know origin is categorical.

One way to avoid the colorbar is to loop over the categories and assign colors based on the category.

See: <https://stackoverflow.com/questions/26139423/plot-different-color-for-different-categorical-levels-using-matplotlib>

```
[14]: colors = {1: 'tab:blue', 2: 'tab:orange', 3: 'tab:green'}
      fig, ax = plt.subplots()
      for key, group in data.groupby(by='origin'):
          group.plot.scatter('mpg', 'model year', c=colors[key], label=key, ax=ax);
```



## 1.5 Seaborn

Seaborn is a Python data visualization library based on matplotlib. It provides a high-level interface for drawing attractive and informative statistical graphics.


<http://seaborn.pydata.org/index.html>

Seaborn is usually imported as `sns`

```
[15]: import seaborn as sns
```


Let's re-create the histograms by origin with seaborn with the figure level `displot()` function.

```
[16]: # Use origin to split mpg into columns
sns.displot(x='mpg', col='origin', data=data);
```



We can display the counts in the same plot, one on top of the other.


```
[17]: # Use origin to color (hue) in the same plot
sns.displot(x='mpg', hue='origin', data=data);
```



To have an idea of the split between origin, we can stack the counts, adding up to total.


```
[18]: sns.displot(x='mpg', hue='origin', data=data, multiple='stack');
```





We can look at the differences in origins with a boxplot too

```
[19]: sns.boxplot(x='origin', y='mpg', data=data);
```




Let's re-create the scatter plot to see if mpg and model year are correlated by origin.

To make the legend show strings we will create a origin string column with strings naming the origins rather than 0 and 1.


```
[20]: data['origin_str'] = data['origin'].replace([1, 2, 3], ['Origin 1', 'Origin 2', 'Origin 3'])
```

```
[21]: ax = sns.scatterplot(x='mpg', y='model year', data=data, hue='origin_str')
```



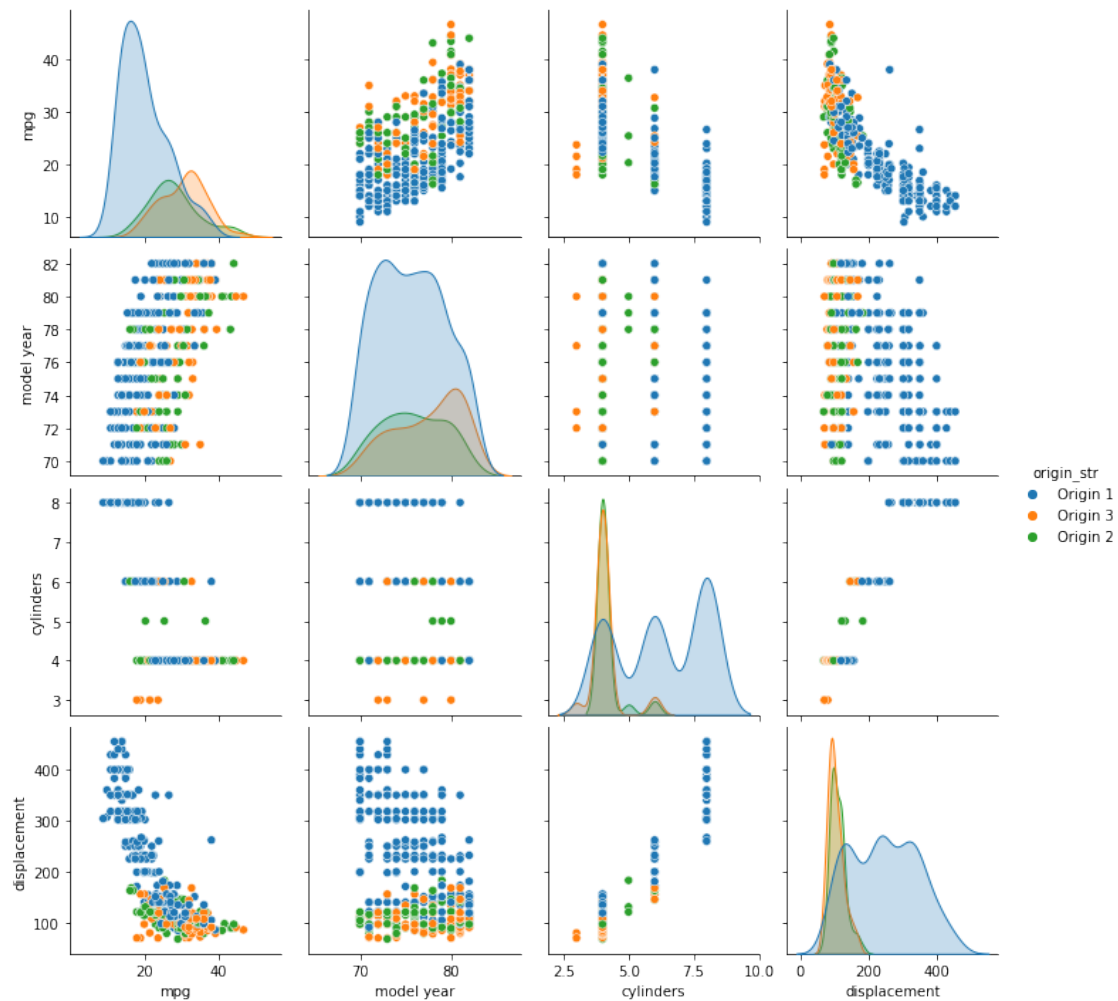
Adding a regression line helps with visualizing the relationship

```
[22]: ax = sns.lmplot(x='mpg', y='model year', data=data, hue='origin_str')
```




Maybe there are other correlations in the data set. Pairplot is a great way to get an overview

```
[23]: sns.pairplot(data, vars=['mpg', 'model year', 'cylinders', 'displacement'],  
    ↪ hue='origin_str');
```



As an alternative, we can visualize the correlation matrix as a heatmap

```
[24]: g = sns.heatmap(data[['mpg', 'model year', 'cylinders', 'displacement']].
    ↪corr(method='spearman'),
    annot=True)
```



There are nice tutorials on the Seaborn website, be sure to check these out.

[ ]: